

Company news component — project report

A standalone backend that keeps a news feed for every company in a **1,515-company catalogue**. Its own process, its own database, no UI.

Every figure below comes from one run and is reproducible with `npm run metrics`. Nothing is estimated.

Results

Companies in catalogue	1,515	Companies returning news	1,163 (76.8%)
Resolved to a real security	1,463 (96.6%)	Attribution certain	78.6% of articles
Exchange listings found	1,987	Full run	26 min
Depository receipts found	140	Refusals / failures	0 / 0

Recurring cost: \$0. No paid API is used anywhere, and no language model is called at any point.

The idea that makes it work

Free news APIs only cover companies listed in America, and a third of this catalogue isn't. But most foreign companies **also** trade in America under a different symbol — Toyota is `7203` in Tokyo and `TM` in New York.

So before fetching any news, the service resolves every ticker to **all** the venues it trades on, including its US depository receipt. That brings **94.6%** of the catalogue within reach of a free, ticker-native provider — which is why no paid API is required.

Segment	Companies	Share
US exchange listing	1,252	82.6%
US OTC line only	182	12.0%
Reachable with a US symbol	1,434	94.6%
No US line at all	29	1.9%
Unresolved	52	3.4%

A ticker is not an identifier

The most important correctness problem in the data. **Seven tickers resolve to a different company than the one named:**

<code>BBY</code>	Balfour Beatty → Best Buy	<code>ADM</code>	Admiral → Archer-Daniels	<code>NOV</code>	Novo Nordisk → NOV Inc
<code>ENR</code>	Siemens Energy → Energizer	<code>FTK</code>	flatexDEGIRO → Flotek	<code>CWK</code>	Cranswick → Cushman
<code>MOVE</code>	Medacta → Corvex				

Filtering by exchange does **not** fix this — `ADM` narrowed to London still returns Archer-Daniels-Midland, because it trades there too. Only checking that the company *name* agrees does.

Untreated, these seven produce a full, plausible, entirely wrong feed. All seven now resolve correctly, and each gained a US symbol in the process.

Why some companies have no news

"All 1,515 fetched correctly" is already true — zero refusals, zero failures. Every company was asked the right source with the right symbol. What varies is whether news existed:

Cause	Companies	Fixable?
Genuinely quiet that week	~272	No — and shouldn't be. A clean zero is a real answer
Delisted or acquired	45	No — <code>EA</code> and <code>EQR</code> no longer exist as securities
Thin OTC-only coverage	~30	Yes, with a paid provider
Renamed beyond recognition	4	Needs a manual alias

Design decisions worth knowing

- Provider-agnostic.** Sources sit behind one interface, ordered by config. Adding one is a config change, not a rewrite.

- **Idempotent.** Articles are keyed on a canonical URL, so a daily re-run stores nothing new. Verified: zero duplicate URLs across 6,461 articles.
- **Wrong attributions are rejected, not hidden.** Name search filed a film review under Kid ASA and every *Jensen Huang* story under Jensen-Group. 3,041 such articles were dropped at ingest in this run.
- **"No news" ≠ "provider refused."** They look identical if you only count articles. The first is normal; the second means coverage is dying. They are recorded separately, along with every provider tried per company — because a working fallback can hide a broken primary.
- **No language model.** Article text is written by strangers; letting a model read it and decide what gets stored would hand outsiders influence over the pipeline. The rules used instead are unit-tested against real failures.

Do we need to pay?

Not for coverage — 94.6% is already served free and ticker-native.

Possibly for provenance. 97% of company news comes from three aggregators (Yahoo, Benzinga, Seeking Alpha); only 2.3% is primary wire. If being early or citing a primary source matters, that is the upgrade worth buying.

Marketaux was evaluated on its free tier: real depth for major international names (4,960 articles for Legal & General), but it does not cover Chinese A-shares and did not beat the free stack on recent-window coverage. **Recommendation: stay free.** The adapter is built and activates on a key, so the decision is reversible.

Known limits

- **52 companies (3.4%) unresolved** — 45 are delisted, which is correct.
- **Name matching is the weakest component.** Some pairs are undecidable from names alone: "Adobe Systems" vs "Adobe" is indistinguishable from "Prudential" vs "Prudential Financial" — the first is one company, the second is two. Those route to a flagged, reduced-confidence path rather than a guess.
- **One ingest worker**; parallel workers need shared rate limiting.
- **A 26-minute run** bounded by the free tier's 60 requests/minute, not by the code. **First start** takes ~9 minutes with a free OpenFIGI key, ~38 without.

Full detail: README.md · COMPARISON.md · API.md · OPERATIONS.md · `data/listings-mapping.csv`