

# Company news component — project report

A standalone backend that keeps a news feed for every company in a **1,515-company catalogue**. Its own process, its own database, no UI.

Every figure comes from one run and is reproducible with `npm run metrics` (coverage gaps: `npm run export:gaps`). Article counts are a point-in-time snapshot of a live corpus — triggering a fetch moves them.

## Results

|                             |               |                            |                |
|-----------------------------|---------------|----------------------------|----------------|
| Companies in catalogue      | 1,515         | Companies returning news   | 1,007 (66.5%)  |
| Resolved to a real security | 1,465 (96.7%) | Attribution certain        | 99.8% of links |
| Exchange listings found     | 1,990         | Articles carrying an image | 81.3%          |
| Depository receipts found   | 140           | Refusals / failures        | 0 / 0          |

**Recurring cost: \$0.** No paid API, and no language model is called anywhere.

### What each row means

- **Resolved to a real security** — for 96.7% of tickers we worked out exactly which real, tradeable share it refers to.
- **Attribution certain** — 99.8% of stored articles came from a provider that *knew the ticker*. Only 0.2% rest on matching a company name.
- **Exchange listings** — those 1,465 companies trade in 1,990 places in total; many trade on more than one exchange.
- **Depository receipts** — 140 of those listings are a foreign share wrapped for trading elsewhere: 139 American (ADR) and 1 global (GDR).
- **Refusals / failures** — no provider refused us, and nothing crashed.

## The idea that makes it work

**The problem:** free news APIs only cover companies listed in America. A third of this list isn't American.

**The trick:** most foreign companies *also* trade in America under a different code. Toyota is 7203 in Tokyo, but also TM in New York. That American version is a **depository receipt** — the same company, wrapped so Americans can buy it.

So before fetching any news, we find every place each company trades. If any of them is American, a free US news API can cover it.

| Segment                           | Companies            | What it means                                                                               |
|-----------------------------------|----------------------|---------------------------------------------------------------------------------------------|
| US exchange listing               | 1,253                | Trades on NYSE or Nasdaq — well covered                                                     |
| US OTC line only                  | 183                  | An American presence, but on the quieter over-the-counter market, where coverage is thinner |
| <b>Reachable with a US symbol</b> | <b>1,436 (94.8%)</b> | The two above: we have <i>some</i> American code to ask about                               |
| No US line at all                 | 29                   | Trades only abroad — must fall back to name search                                          |
| Unresolved                        | 50                   | Ticker matched no identifier in our sources                                                 |

**94.8% is the whole argument.** Because the listing work was done first, almost everything became reachable with free tools. Without it, we would have had to buy a global news service.

## A ticker is not an identifier

A ticker is just a short code **on one particular exchange**. The same code means different companies on different exchanges:

- BBY in **London** = Balfour Beatty, a construction firm
- BBY in **New York** = Best Buy, an electronics retailer

Ask for "news for BBY" and you get **Best Buy's news filed under Balfour Beatty**. The feed looks completely normal — full of real articles — and is about the wrong company entirely. That is what makes it dangerous: nothing looks broken.

**We found seven of these:**

|                                      |                                     |                                   |                                       |
|--------------------------------------|-------------------------------------|-----------------------------------|---------------------------------------|
| BBY Balfour Beatty → <b>Best Buy</b> | ADM Admiral → <b>Archer-Daniels</b> | NOV Novo Nordisk → <b>NOV Inc</b> | ENR Siemens Energy → <b>Energizer</b> |
| FTK flatexDEGIRO → <b>Flotek</b>     | CWK Cranswick → <b>Cushman</b>      | MOVE Medacta → <b>Corvex</b>      |                                       |

**Filtering by exchange does not fix it.** The obvious fix is to ask for "BBY, *in London*". That fails too, because Archer-Daniels-Midland ( ADM in New York) **also trades in London** — so "ADM in London" still returns the wrong company.

**The only reliable check is the name.** We ask the provider what company it thinks the ticker is, then compare that to the name in our list. If they disagree, we reject it.

All seven now resolve correctly — and each **gained a US symbol in the process**. While correctly identifying Balfour Beatty we also found its American code, BAFBF, so its news can now come from a free US provider.

## Wrong articles are thrown away, not hidden

For companies we can look up **by ticker**, the provider tells us exactly which company a story is about — no guessing. But for companies with no US listing there is no ticker to ask with, so we fall back to **searching by company name** — and a name search cannot tell a company called Jensen from a person called Jensen.

**Jensen-Group** is a Belgian maker of industrial laundry machines; its name shortens to "Jensen". **Jensen Huang** is the CEO of Nvidia, one of the most written-about people in business. Searching for "Jensen" returned hundreds of Nvidia stories, filed under a small Belgian laundry-equipment maker.

Same pattern with **Kid ASA**, a Norwegian home-textiles retailer. Its name shortens to "Kid"; searching returned a review of a film called *Club Kid*.

**What we do about it.** Before saving any article found by name search, we check it is genuinely about that company:

1. **Reduce the name to its distinctive words** — Jensen-Group → jensen.
2. **Multi-word names need two words present.** "Munich Reinsurance lifts profit guidance" keeps; "Munich hosts Oktoberfest" is rejected.
3. **Single-word names need the word *and* business vocabulary** — shares, earnings, revenue, results, acquisition, dividend. "Jensen Huang could redefine NVIDIA" is rejected; "Jensen Group reports record revenue" keeps.

Articles that fail are **never stored** — not saved and quietly marked "low quality". **3,041 were thrown away in the last run.** A wrong article never reaches the feed at all.

## Which companies have no news, and why

The full list is exported as data — **data/coverage-gaps.csv**, one row per company with a gap\_reason, regenerated with `npm run export:gaps`. **508 of 1,515** produced no news. They are not one problem:

| Gap reason                       | Companies | Whose problem is it?                                                                              |
|----------------------------------|-----------|---------------------------------------------------------------------------------------------------|
| no_news_us_exchange              | 276       | <b>Nobody's.</b> Identified correctly, asked correctly, genuinely quiet that week                 |
| no_news_otc_only                 | 154       | Thin OTC coverage — a paid provider would help                                                    |
| no_news_no_us_line               | 29        | Trades only abroad; no free US-centric source reaches it                                          |
| unresolved_frankfurt_line        | 17        | Ours — .F Frankfurt lines are not in our identifier sources                                       |
| unresolved_london_secondary_line | 16        | Ours — same, for 0XXX LSE codes                                                                   |
| unresolved_absent_from_directory | 11        | Absent from Finnhub's US directory                                                                |
| unresolved_malformed_ticker      | 3         | The catalogue's — LVMH.F uses _ where every other row uses .                                      |
| unresolved_name_check_rejected   | 1         | The safety check correctly refusing a ticker collision                                            |
| duplicate_listing_line           | 1         | <b>Nothing is missing</b> — same company already covered, <i>with news</i> , under another ticker |

**"Unresolved" means one specific thing:** the ticker could not be matched to an identifier in the sources we use. It does **not** mean the company is delisted. Checking them individually, most are actively traded:

- **16 are London 0XXX lines** — 0FIN Orkla, 0IU8 Safran, 0NQM Vinci, 0R22 Barrick Gold. Large, actively traded companies. The 0XXX code is an LSE convention our identifier sources do not index.
- **17 are Frankfurt .F lines** — EADS.F Airbus, MAK.S.F Marks & Spencer, CMXH.F CSL. Same story.
- **1 is a duplicate that costs nothing** — NEMC.L Newmont is already covered as 0R28, which holds 12 articles. gap\_reason counts a row as a duplicate only when its twin actually *has news*; six other same-name pairs exist, but both sides are empty, so

those are counted as real gaps rather than written off.

- **11 are plain US tickers absent from Finnhub's 30,995-symbol US directory** — EA , SKX , EQR , JHG . Some are genuine take-privates; others look like directory gaps. We do not claim to know which without checking each.

**The name check now resolves trading names and native-language names.** A surface-form comparison cannot know that "Westinghouse Air Brake Technologies" and "WABTEC CORP" are one company, or that "MUENCHENER RUECKVER AG-REG" is Munich Re in German — both were correct identifiers rejected for looking different. A small, explicit alias table handles them, and both now resolve.

The one remaining rejection is **correct and deliberate**: P matches "EVERPURE INC-A" in the US directory, while the catalogue row is Pure Storage. Those are different companies, and refusing the match is the same protection that caught the seven ticker collisions above.

**The catalogue also double-counts.** 1,515 rows are not 1,515 distinct businesses: Novartis appears three times ( NOVN , NOVNEE , 0QLR ), SAP three times ( SAP , 0NW4 , SAPG.F ), Alphabet twice ( G00GL , 0HD6 ). Coverage percentages are therefore computed per row, and the true per-business figure is somewhat better than the headline suggests.

## What is actually fixable

| Fix                                                                          | Companies | Status                               |
|------------------------------------------------------------------------------|-----------|--------------------------------------|
| Alias table for trading names and non-English names                          | 2         | <b>Done</b> — WAB , MUV2 now resolve |
| Accept <code>_</code> as a venue separator alongside <code>.</code>          | 3         | <b>Done</b> — applied at lookup      |
| Map <code>0XXX / .F</code> codes via an LSE/Deutsche Börse identifier source | 33        | Needs a second identifier source     |
| Paid provider for OTC and non-US lines                                       | up to 183 | Cost, not code                       |

The first two are done and reflected in the figures above. The remaining two are not code problems: one needs an identifier source that indexes European secondary lines, the other needs a licence.

## Every link now points at the publisher

Two separate faults made stored links fail when clicked. Both are fixed.

**1. Google News redirects (~1,400 articles).** These arrived as `news.google.com` links rather than publisher addresses, and cannot be repaired: the token decodes to an opaque Google identifier rather than the article's address, the page returns nothing to a non-browser client, and browsers frequently refuse the redirect. Such a story can never show a picture either. We stopped storing them.

**2. Finnhub redirect wrappers (4,975 articles — 98% of the feed).** Finnhub does not return the article's URL. It returns a link into *its own* domain, `finnhub.io/api/news?id=...`, which 302s onward to the publisher. Storing the wrapper rather than resolving it costs three things:

- every reader paid an extra hop, and the link dies entirely if Finnhub is down — a news archive should not depend on the liveness of the API it came from;
- the wrapper **hid the destination**, so a dead publisher could not be filtered: the URL looked healthy right up until the click;
- **dedupe degraded**, because every wrapper carries a distinct `id`, so one wire story syndicated to three outlets was three separate articles.

Ingest now resolves the wrapper to the publisher's own address before storing, and a backfill rewrote all 4,975 existing rows. **0 wrappers remain**, across **77 distinct publisher domains**.

**What the wrapper was hiding: chartmill.com.** With destinations visible, one host turned out to be dead — it refuses the connection outright, returning no HTTP status at all in ~0.35s, with a browser user-agent, which a reader sees as a 504. It accounted for 189 articles and 188 broken images. It is now filtered, via a configurable `DEAD_ARTICLE_HOSTS` list.

That list is deliberately tiny and evidence-based. It is **not** a list of sites that block automated clients: Benzinga and Seeking Alpha both refuse our requests while serving browsers perfectly well, and dropping them would discard 2,298 working articles. A host qualifies only when it fails for a real browser too.

|                              | Before      | After |
|------------------------------|-------------|-------|
| Links via a redirect wrapper | 4,975 (98%) | 0     |
| Links on dead hosts          | 189         | 0     |
| Broken Google links          | ~1,400      | 0     |
| Articles carrying an image   | 65%         | 81.3% |
| Companies returning news     | 1,163       | 1,007 |

The lost coverage is real, and it is the clearest argument for a paid provider: one returns genuine publisher links and images for exactly these companies.

## What happens when a reader clicks

Links now go straight to the publisher. That fixed the broken ones — but it exposes a second problem that is worth stating plainly, because it is the clearest argument for buying content.

**Every headline leaves our platform**, and where it lands is the publisher's site, not ours:

| Publisher                         | Share of articles | What the reader sees                  |
|-----------------------------------|-------------------|---------------------------------------|
| <b>benzinga.com</b>               | <b>33.4%</b>      | Often a bot check or a sign-up prompt |
| finance.yahoo.com                 | 30.8%             | Opens normally                        |
| <b>seekingalpha.com</b>           | <b>15.8%</b>      | Often a registration wall             |
| fool.com, cnbc.com, and 72 others | ~20%              | Mostly opens normally                 |

**So roughly half of all clicks (49%) land on a site that may ask the reader to register before showing the article.** Nothing is broken — these are the correct publisher URLs, correctly resolved. The reader simply meets someone else's paywall on a link we presented as ours, and a different site layout each time.

**We did not remove them.** Benzinga and Seeking Alpha serve real browsers perfectly well; they only refuse automated clients, so we cannot reliably detect a wall in advance. Dropping both would delete **49% of the catalogue's articles** to avoid an inconvenience that many readers never hit.

**This is not fixable with the current providers.** Finnhub and Marketaux license a headline, a teaser and a **link** — their terms require the link out. Rendering the full text anyway, or scraping the publisher's page for it, is republishing copyrighted work. The fix is to license content we are permitted to display, which is what the recommendation below is about.

## The component in use

The service has no UI of its own — it is a backend. These are from the existing research app consuming it over its HTTP API, which is how it would be integrated.

🔍 Search for stocks, ETFs, and more

My StocksNewsCatalogueGlobal NewsAlerts27My StrategySettings🌙👤

Global News

The full production catalogue, served by the standalone news component over its HTTP API.

Companies tracked1,515

Resolved to a listing1,463

Exchange listings1,987

Articles stored4,882

Companies (1,515)

Search name or tickExchangeSearch

Shenzhen Airport Co Ltd

000089 · CN · hint SZSE

C2:000089

0no news

Midea Group Co Ltd

000333 · CN · hint SZSE

C2:000333US:MQCLY ADR

0no news

Luzhou Lao Jiao Co Ltd

000568 · CN · hint SZSE

C2:000568

0no news

Sectra AB

0A0L · GB · hint LSE

LN:0A0LUS:STKTRUS:SCTAY ADR

0no news

Fox Corporation

0A0X · GB · hint LSE

LN:0A0XUS:FOXAUS:FOX

23d ago

Uber Technologies Inc

0A1U · GB · hint LSE

LN:0A1UUS:UBER

342d ago

Bloom Energy Corporation

0A4L · GB · hint LSE

LN:0A4LUS:BE

241d ago

DexCom, Inc.

0A4M · GB · hint LSE

LN:0A4MUS:DXCM

32d ago

ACADIA Pharmaceuticals Inc.

0A4W · GB · hint LSE

LN:0A4WUS:ACAD

43d ago

Jack Henry & Associates, Inc.

0A6D · GB · hint LSE

LN:0A6DUS:JKHY

153d ago

The Cigna Group

0A77 · GB · hint LSE

LN:0A77US:CI

62d ago

Carvana Co.

0A79 · GB · hint LSE

LN:0A79US:CVNA

43d ago

Palantir Technologies Inc.

GB · hint LSE

LN:0A7RUS:PLTR

442d ago

Market & macro

Stories that move prices without naming a company — conflict, rates, disasters, policy. From primary wires.

CP Capital Acquires Ascent Athens, a 200-Unit Class A Multifamily Community in Athens, Georgia

BusinessWire · 24h ago

aggregator

Family offices are making a bullish bet on the stock market, according to CNBC Family Office Portfolio Tracker

CNBC · 24h ago

primary wire

This is 'one of the worst' credit card habits, experts say—41% of cardholders do it

CNBC · 1d ago

primary wire

First Financial Corporation and First Illinois Corporation Sign Merger Agreement

GlobalNewsWire · 1d ago

aggregator

Kansas City Fed's Schmid says inflation 'stubborn' and 'sticky,' policy rate not restrictive

CNBC · 1d ago

primary wire

Form 8.3 Picton Property Income Limited

GlobalNewsWire · 1d ago

aggregator

Company news

Nvidia's Unusual Move Answers AI Trade's Key Question; \$40 Trillion US Debt Bomb Ticks; Key Fed Speech Awaited

AAPL · Benzinga · 1d ago

aggregator

ETJ: Very High Tech Exposure Could Be A Risk

AAPL · SeekingAlpha · 1d ago

aggregator

Dollar Tree: The Margin Of Safety Is Finally Gone (Downgrade)

OIC8 · SeekingAlpha · 1d ago

aggregator

VICI Properties: When It Rains Gold, Put Out The Bucket

VICI · SeekingAlpha · 1d ago

aggregator

Realty Income: Business Expansion Has Not Accelerated Growth

· SeekingAlpha · 1d ago

aggregator

KKR & Co. Inc. Q2: 25% Capital Appreciation Potential From Its Baby Bonds

KKR · SeekingAlpha · 1d ago

aggregator

Research

Hi Insar, ask any financial question

Give me an overview of the market

🔍

Why did the company stock move after its latest earnings...

🔍

What was the company's revenue and EPS for the most...

🔍

Explore what's possible

+ Summarise my portfolio

🔍 Check my strategy fit

🔍 Deep search the news

🔍 Analyse my watchlist

Ask anything

+🔍↑

Every company in the catalogue, with the venues it resolved to. LN:0A1U is the London line, US:UBER the US symbol news is actually fetched with. A 0 with "no news" is a clean zero, not an error — the distinction the run reporting is built around.

Search for stocks, ETFs, and more

My StocksNewsCatalogueGlobal NewsAlerts27My StrategySettings

Global News

The full production catalogue, served by the standalone news component over its HTTP API.

Companies tracked1,515

Resolved to a listing1,463

Exchange listings1,987

Articles stored4,882

Companies (1)

MicrosoftExchangeSearchClear

Microsoft Corp

MSFT · US · no exchange hint

US:MSFT

1172d ago

Market & macro

Stories that move prices without naming a company — conflict, rates, disasters, policy. From primary wires.

CP Capital Acquires Ascent Athens, a 200-Unit Class A Multifamily Community in Athens, Georgia

BusinessWire · 24h ago

aggregator

Family offices are making a bullish bet on the stock market, according to CNBC Family Office Portfolio Tracker

CNBC · 24h ago

primary wire

This is 'one of the worst' credit card habits, experts say—41% of cardholders do it

CNBC · 1d ago

primary wire

First Financial Corporation and First Illinois Corporation Sign Merger Agreement

GlobalNewsWire · 1d ago

aggregator

Kansas City Fed's Schmid says inflation 'stubborn' and 'sticky,' policy rate not restrictive

CNBC · 1d ago

primary wire

Form 8.3 Picton Property Income Limited

GlobalNewsWire · 1d ago

aggregator

Research

Hi Insar, ask any financial question

Give me an overview of the market

Why did the company stock move after its latest earnings...

What was the company's revenue and EPS for the most...

Explore what's possible

+ Summarise my portfolio

Check my strategy fit

Deep search the news

Analyse my watchlist

Ask anything

+🔊⬆

Search by name or ticker. The result shows the resolved listing (US:MSFT) and the stored article count.

Search for stocks, ETFs, and more

My StocksNewsCatalogueGlobal NewsAlerts27My StrategySettings

← Back to Global News

Microsoft Corp

MSFT · US · resolved

US:MSFT

News (57)

GTA 6 Leaker Cashes Out \$250,000 Meme Coin Hours Before Rockstar's Reveal

aggregator

GTA 6 leaker cashes out over \$250K from a self-created Solana memecoin hours before Rockstar's gameplay reveal.

Benzinga · 1d ago

What's Going On With Microsoft Stock?

aggregator

Microsoft Corp (NASDAQ:MSFT) shares are rising Thursday. The company announced a new AI partnership with Saudi Arabia's HUMAIN, with broader strength across software stocks adding to the momentum.

Benzinga · 1d ago · also mentions 1 other

Datadog: Growing Into A 400x P/E Will Be Difficult

aggregator

Datadog shows impressive 40%+ revenue CAGR, but trades at over 400x P/E and a PEG above 3. Click to read why DDOG is rated a Sell.

SeekingAlpha · 1d ago · also mentions 1 other

Bill Gates Wants to Hit the Brakes. Nvidia Is Financing AI's Next Sprint.

aggregator

Bill Gates wants AI to slow down. Nvidia's latest earnings reveal how it's financing AI infrastructure to keep the boom accelerating.

Benzinga · 1d ago · also mentions 1 other

ChronoScale Announces Plans With Microsoft For 50-MW AI Compute Deployment In North America

aggregator

ChronoScale Corporation (NASDAQ:CHRN), (634;ChronoScale34; or the 634;Company34;), an accelerated compute platform purpose-built to support demanding artificial intelligence workloads, today announced plans with

Benzinga · 1d ago

Trump Weighs New Semiconductor Tariffs That Could Hit AI Data Center Servers, Laptops and Gaming Consoles: Report

aggregator

Trump weighs broader semiconductor tariffs as tech giants warn higher costs could threaten U.S. AI data center expansion.

Benzinga · 1d ago · also mentions 2 other

Understanding Microsoft's Position In Software Industry Compared To Competitors

aggregator

In the dynamic and cutthroat world of business, conducting thorough company analysis is essential for investors and industry experts. In this article, we will undertake a comprehensive industry comparison, evaluating

Benzinga · 1d ago

NVIDIA CFO Says We Have Invested \$50B In Frontier Labs; OpenAI's Existing, Planned Commitments Represent Around 12GW Of NVIDIA Compute

aggregator

-Conference Call

Benzinga · 2d ago · also mentions 1 other

IREN Q4 Preview: Stock Down 48% From Highs, Will Data Center Monetization Shine Through In

aggregator

Research

Hi Insar, ask any financial question

Give me an overview of the market

Why did the company stock move after its latest earnings...

What was the company's revenue and EPS for the most...

Explore what's possible

+ Summarise my portfolio

Check my strategy fit

Deep search the news

Analyse my watchlist

Ask anything

+🔊⬆

One company's feed. "also mentions 1 other" is the shared-article model: a story naming several companies is stored once and linked to each, never duplicated per company.

**Round-ups are filtered out of a company's feed.** Aggregators tag a single macro or sector article against every ticker it mentions — "10 IT stocks with whale alerts" arrives filed against ten companies at once. Attribution is *certain* (the provider named the ticker) but the story is not *about* any one of them, and left unfiltered those dominate the feed of a widely-mentioned company like Microsoft.

A company's own feed therefore defaults to `max_companies=3` : an article filed against more than three companies is a list, not company news. The relevance score also carries a round-up penalty applied at read time, where the live company count is known. Callers that want the unfiltered set pass `max_companies` explicitly.

## The two timing numbers

They measure different things.

**A 26-minute run** — fetching news for every company, day to day. Finnhub's free plan allows 60 requests per minute, and we ask about one company per request.  $1,434 \text{ companies} \div 55 \text{ per minute} \approx 26 \text{ minutes}$ . **The code is not slow — it is deliberately waiting** to stay inside the free limit. A paid tier would cut this to a few minutes and change nothing else.

**A ~9-minute first start** — this happens **once**. On first boot we look up all 1,515 tickers to find their exchanges, via OpenFIGI. Without a key that allows 25 lookups a minute; with a **free** key, 250 a minute and larger batches — hence 9 minutes rather than 38. After the first time it never happens again.

## Design decisions worth knowing

- **Provider-agnostic.** Sources sit behind one interface, ordered by config. Adding one is a config change, not a rewrite.
- **Idempotent.** Articles are keyed on a cleaned-up URL, so a daily re-run stores nothing new. Verified: zero duplicate URLs.
- **"No news" is not "provider refused."** They look identical if you only count articles. The first is normal; the second means coverage is dying. They are recorded separately, along with every provider tried per company — because a working fallback can hide a broken primary.
- **No language model.** Article text is written by strangers; letting a model read it and decide what gets stored would hand outsiders influence over the pipeline. The rules used instead are unit-tested against real failures.

## Do we need to pay? — and what to buy

**Not for coverage.** 94.8% is already reachable free and ticker-native.

**Yes, for the reading experience.** Three problems are not solvable with any free tier, and they are the ones a user actually feels:

1. **Every headline leaves our platform.** The reader lands on Yahoo, Benzinga, Seeking Alpha or Fool — a different site, a different layout, each time.
2. **Some of those sites demand a sign-up or a subscription** before showing the article. We cannot tell in advance which, so the reader plays paywall roulette on links we presented as ours.
3. **We cannot legally fix that ourselves.** Free and low-cost APIs — Finnhub, Marketaux, Alpha Vantage — license a **headline, a short teaser and a link**. Their terms require the link out. Displaying the full text we did not license, or scraping the publisher's page for it, is republishing someone else's copyrighted work. Financial publishers enforce this actively, and it is not a risk worth taking to save a click.

**So the requirement "show the article in our own UI" is a licensing decision, not an engineering one.** The fix is to buy content we are permitted to display.

## Recommendation: Benzinga's licensed newsfeed as primary

Benzinga owns its newsroom, so it can license what an aggregator cannot. Its paid tiers are **explicitly embeddable** — the full body and the image may be published inside our own interface, with no redirect and no third-party sign-in. Its free tier is headline + teaser + link, which is exactly the constraint we are trying to escape.

| Requirement                          | How Benzinga meets it                                                                                         |
|--------------------------------------|---------------------------------------------------------------------------------------------------------------|
| Show the article without redirecting | Paid tiers license the <b>full body and image</b> for display on our platform                                 |
| Always fresh, automatically          | Pull REST with <code>updatedSince</code> , plus <b>WebSocket, TCP stream and webhooks</b> — push, not polling |
| No per-company rate ceiling          | A firehose of the newsroom, not one request per company — the cost stops scaling with catalogue size          |
| One consistent source                | One publisher, one voice, one layout — no paywall roulette                                                    |



That last row also solves the scale problem in the section above: polling costs one request per company per day, so 50,000 companies is 50,000 requests. A stream costs the same regardless of how many companies we track.

**The honest limitation.** Benzinga publishes roughly 130–160 full articles a day from a US-focused newsroom. That is deep, not wide — it will not cover Chinese A-shares or thin European lines. It replaces the *quality* of our feed, not all of its *breadth*.

## The two-provider architecture

Exactly two, which is the practical minimum that still covers the catalogue:

| Tier     | Provider        | Licence              | Reader experience                                                |
|----------|-----------------|----------------------|------------------------------------------------------------------|
| Primary  | Benzinga (paid) | Full body embeddable | Reads in our UI. No redirect, no sign-in                         |
| Fallback | Marketaux       | Headline + link      | Opens the publisher — only for companies Benzinga does not cover |

Everything needed to make this switch already exists: providers sit behind one adapter interface and are ordered by `NEWS_PROVIDER_ORDER`, so adding Benzinga is a new adapter plus a config change, and Finnhub can be dropped the day the licensed feed proves out. The article schema already carries `image_url` and a summary; a licensed body needs one nullable `body` column and a UI that renders it instead of linking out.

**Pricing is not published** — content licensing is quoted per customer, via `licensing@benzinga.com`. The brief offers to provision a paid key, so the ask is a quote for the embeddable-body tier at ~1,500 companies, with room to grow.

## Known limits

- **50 companies (3.3%) unresolved** — see the table above; most are foreign secondary listings our identifier sources do not index, not delistings.
- **Name matching is the weakest component.** Some pairs are undecidable from names alone: "Adobe Systems" vs "Adobe" is indistinguishable from "Prudential" vs "Prudential Financial" — the first is one company, the second is two. Those route to a flagged, reduced-confidence path rather than a guess.
- **Some publishers block us.** Bloomberg returns an error to automated requests, so those stories carry no image. We do not disguise the service as a browser to get around it.
- **One ingest worker**; parallel workers need shared rate limiting.

Full detail: README.md · COMPARISON.md · API.md · OPERATIONS.md · `data/listings-mapping.csv`