

Company news component — project report

A standalone backend that keeps a news feed for every company in a **1,515-company catalogue**. Its own process, its own database, no UI.

Every figure comes from one run and is reproducible with `npm run metrics` (coverage gaps: `npm run export:gaps`). Article counts are a point-in-time snapshot of a live corpus — triggering a fetch moves them.

Results

Companies in catalogue	1,515	Companies returning news	1,005 (66.3%)
Resolved to a real security	1,463 (96.6%)	Attribution certain	99.8% of links
Exchange listings found	1,987	Articles carrying an image	82.0%
Depository receipts found	140	Refusals / failures	0 / 0

Recurring cost: \$0. No paid API, and no language model is called anywhere.

What each row means

- **Resolved to a real security** — for 96.6% of tickers we worked out exactly which real, tradeable share it refers to.
- **Attribution certain** — 99.8% of stored articles came from a provider that *knew the ticker*. Only 0.2% rest on matching a company name.
- **Exchange listings** — those 1,463 companies trade in 1,987 places in total; many trade on more than one exchange.
- **Depository receipts** — 140 of those listings are a foreign share wrapped for trading elsewhere: 139 American (ADR) and 1 global (GDR).
- **Refusals / failures** — no provider refused us, and nothing crashed.

The idea that makes it work

The problem: free news APIs only cover companies listed in America. A third of this list isn't American.

The trick: most foreign companies *also* trade in America under a different code. Toyota is 7203 in Tokyo, but also TM in New York. That American version is a **depository receipt** — the same company, wrapped so Americans can buy it.

So before fetching any news, we find every place each company trades. If any of them is American, a free US news API can cover it.

Segment	Companies	What it means
US exchange listing	1,252	Trades on NYSE or Nasdaq — well covered
US OTC line only	182	An American presence, but on the quieter over-the-counter market, where coverage is thinner
Reachable with a US symbol	1,434 (94.6%)	The two above: we have <i>some</i> American code to ask about
No US line at all	29	Trades only abroad — must fall back to name search
Unresolved	52	Could not identify the company (mostly delisted)

94.6% is the whole argument. Because the listing work was done first, almost everything became reachable with free tools. Without it, we would have had to buy a global news service.

A ticker is not an identifier

A ticker is just a short code **on one particular exchange**. The same code means different companies on different exchanges:

- BBY in **London** = Balfour Beatty, a construction firm
- BBY in **New York** = Best Buy, an electronics retailer

Ask for "news for BBY" and you get **Best Buy's news filed under Balfour Beatty**. The feed looks completely normal — full of real articles — and is about the wrong company entirely. That is what makes it dangerous: nothing looks broken.

We found seven of these:

BBY Balfour Beatty → Best Buy	ADM Admiral → Archer-Daniels	NOV Novo Nordisk → NOV Inc	ENR Siemens Energy → Energizer
FTK flatexDEGIRO → Flotek	CWK Cranswick → Cushman	MOVE Medacta → Corvex	

Filtering by exchange does not fix it. The obvious fix is to ask for "BBY, *in London*". That fails too, because Archer-Daniels-Midland (ADM in New York) **also trades in London** — so "ADM in London" still returns the wrong company.

The only reliable check is the name. We ask the provider what company it thinks the ticker is, then compare that to the name in our list. If they disagree, we reject it.

All seven now resolve correctly — and each **gained a US symbol in the process**. While correctly identifying Balfour Beatty we also found its American code, BAFBF , so its news can now come from a free US provider.

Wrong articles are thrown away, not hidden

For companies we can look up **by ticker**, the provider tells us exactly which company a story is about — no guessing. But for companies with no US listing there is no ticker to ask with, so we fall back to **searching by company name** — and a name search cannot tell a company called Jensen from a person called Jensen.

Jensen-Group is a Belgian maker of industrial laundry machines; its name shortens to "Jensen". **Jensen Huang** is the CEO of Nvidia, one of the most written-about people in business. Searching for "Jensen" returned hundreds of Nvidia stories, filed under a small Belgian laundry-equipment maker.

Same pattern with **Kid ASA**, a Norwegian home-textiles retailer. Its name shortens to "Kid"; searching returned a review of a film called *Club Kid*.

What we do about it. Before saving any article found by name search, we check it is genuinely about that company:

- 1. Reduce the name to its distinctive words** — Jensen-Group → jensen .
- 2. Multi-word names need two words present.** "Munich Reinsurance lifts profit guidance" keeps; "Munich hosts Oktoberfest" is rejected.
- 3. Single-word names need the word *and* business vocabulary** — shares, earnings, revenue, results, acquisition, dividend. "Jensen Huang could redefine NVIDIA" is rejected; "Jensen Group reports record revenue" keeps.

Articles that fail are **never stored** — not saved and quietly marked "low quality". **3,041 were thrown away in the last run.** A wrong article never reaches the feed at all.

Which companies have no news, and why

The full list is exported as data — data/coverage-gaps.csv , one row per company with a gap_reason , regenerated with npm run export:gaps . **510 of 1,515** produced no news. They are not one problem:

Gap reason	Companies	Whose problem is it?
no_news_us_exchange	276	Nobody's. Identified correctly, asked correctly, genuinely quiet that week
no_news_otc_only	154	Thin OTC coverage — a paid provider would help
no_news_no_us_line	29	Trades only abroad; no free US-centric source reaches it
unresolved_frankfurt_line	17	Ours — .F Frankfurt lines are not in our identifier sources
unresolved_london_secondary_line	16	Ours — same, for 0XXX LSE codes
unresolved_absent_from_directory	11	Absent from Finnhub's US directory
unresolved_malformed_ticker	3	The catalogue's — LVMH_F uses _ where every other row uses .
unresolved_name_check_rejected	3	Ours — the safety check fired, twice wrongly
duplicate_listing_line	1	Nothing is missing — same company already covered, <i>with news</i> , under another ticker

A correction worth stating plainly. An earlier draft of this report said the 52 unresolved were "45 delisted, which is correct." That was overconfident. What is actually true is narrower: **45 could not be matched to an identifier in the sources we use.** Checking them individually shows most are not delisted at all:

- 16 are London 0XXX lines** — 0FIN Orkla, 0IU8 Safran, 0NQM Vinci, 0R22 Barrick Gold. Large, actively traded companies. The 0XXX code is an LSE convention our identifier sources do not index.
- 17 are Frankfurt .F lines** — EADS.F Airbus, MAK.S.F Marks & Spencer, CMXH.F CSL. Same story.

- **Only 1 is a harmless duplicate** — `NEMC.L` Newmont, already covered as `ØR28`, which has 12 stored articles. An earlier draft of this report claimed seven such rows and said "nothing is missing for these at all." That was wrong: the other six map to twins that are themselves empty, so the news really is missing. `gap_reason` now requires the twin to *have news* before a row is written off as a duplicate.
- **11 are plain US tickers absent from Finnhub's 30,995-symbol US directory** — `EA`, `SKX`, `EQR`, `JHG`. Some are genuine take-privates; others look like directory gaps. We do not claim to know which without checking each.

Two of the three name-check rejections are our bug, not a bad match. `WAB` was rejected because the catalogue says "Westinghouse Air Brake Technologies" and the directory says "WABTEC CORP" — the same company under its trading name. `MUV2` was rejected because the directory carries Munich Re's German name, "MUENCHENER RUECKVER AG-REG". The matcher compares surface forms and has no notion of an abbreviation or a translation. The third rejection, `P` → "EVERPURE INC-A" against a catalogue row for Pure Storage, is **correct** — those really are different companies, and it is the same protection that caught the seven ticker collisions.

The catalogue also double-counts. 1,515 rows are not 1,515 distinct businesses: Novartis appears three times (`NOVN`, `NOVNEE`, `ØQLR`), SAP three times (`SAP`, `ØNW4`, `SAPG.F`), Alphabet twice (`GØØGL`, `ØHD6`). Coverage percentages are therefore computed per *row*, and the true per-business figure is somewhat better than the headline suggests.

What is actually fixable

Fix	Companies recovered	Effort
Treat <code>_</code> as <code>.</code> when parsing tickers	3	Trivial
Alias table for trade names and non-English names	2	Small, manual
Collapse duplicate listing rows onto one company	1	Small
Map <code>ØXXX / .F</code> codes via an LSE/Deutsche Börse identifier source	29	Medium
Paid provider for OTC and non-US lines	up to 183	Cost, not code

The first three are ~6 companies for an afternoon's work. The last row is the only one that requires spending money, and it is the largest by far.

Every link now points at the publisher

Two separate faults made stored links fail when clicked. Both are fixed.

1. Google News redirects (~1,400 articles). These arrived as `news.google.com` links rather than publisher addresses, and cannot be repaired: the token decodes to an opaque Google identifier rather than the article's address, the page returns nothing to a non-browser client, and browsers frequently refuse the redirect. Such a story can never show a picture either. We stopped storing them.

2. Finnhub redirect wrappers (4,975 articles — 98% of the feed). Finnhub does not return the article's URL. It returns a link into *its own* domain, `finnhub.io/api/news?id=...`, which 302s onward to the publisher. Storing that was a mistake on three counts:

- every reader paid an extra hop, and the link dies entirely if Finnhub is down — a news archive should not depend on the liveness of the API it came from;
- the wrapper **hid the destination**, so a dead publisher could not be filtered: the URL looked healthy right up until the click;
- **dedupe degraded**, because every wrapper carries a distinct `id`, so one wire story syndicated to three outlets was three separate articles.

Ingest now resolves the wrapper to the publisher's own address before storing, and a backfill rewrote all 4,975 existing rows. **0 wrappers remain**, across **76 distinct publisher domains**.

What the wrapper was hiding: `chartmill.com`. With destinations visible, one host turned out to be dead — it refuses the connection outright, returning no HTTP status at all in ~0.35s, with a browser user-agent, which a reader sees as a 504. It accounted for 189 articles and 188 broken images. It is now filtered, via a configurable `DEAD_ARTICLE_HOSTS` list.

That list is deliberately tiny and evidence-based. It is **not** a list of sites that block automated clients: Benzinga and Seeking Alpha both refuse our requests while serving browsers perfectly well, and dropping them would discard 2,298 working articles. A host qualifies only when it fails for a real browser too.

	Before	After
Links via a redirect wrapper	4,975 (98%)	0
Links on dead hosts	189	0
Broken Google links	~1,400	0
Articles carrying an image	65%	82.0%
Companies returning news	1,163	1,005

The lost coverage is real, and it is the clearest argument for a paid provider: one returns genuine publisher links and images for exactly these companies.

What happens when a reader clicks

Links now go straight to the publisher. That fixed the broken ones — but it exposes a second problem that is worth stating plainly, because it is the clearest argument for buying content.

Every headline leaves our platform, and where it lands is the publisher's site, not ours:

Publisher	Share of articles	What the reader sees
finance.yahoo.com	32.2%	Opens normally
benzinga.com	31.6%	Often a bot check or a sign-up prompt
seekingalpha.com	15.6%	Often a registration wall
fool.com, cnbc.com, and 71 others	~20%	Mostly opens normally

So roughly half of all clicks (47%) land on a site that may ask the reader to register before showing the article. Nothing is broken — these are the correct publisher URLs, correctly resolved. The reader simply meets someone else's paywall on a link we presented as ours, and a different site layout each time.

We did not remove them. Benzinga and Seeking Alpha serve real browsers perfectly well; they only refuse automated clients, so we cannot reliably detect a wall in advance. Dropping both would delete **47% of the catalogue's articles** to avoid an inconvenience that many readers never hit.

This is not fixable with the current providers. Finnhub and Marketaux license a headline, a teaser and a **link** — their terms require the link out. Rendering the full text anyway, or scraping the publisher's page for it, is republishing copyrighted work. The fix is to license content we are permitted to display, which is what the recommendation below is about.

The component in use

The service has no UI of its own — it is a backend. These are from the existing research app consuming it over its HTTP API, which is how it would be integrated.

🔍 Search for stocks, ETFs, and more

🔊

My Stocks

News

Catalogue

Global News

Alerts 27

My Strategy

Settings

🌙

👤

Global News

The full production catalogue, served by the standalone news component over its HTTP API.

last fetch 2h ago - succeeded

Companies tracked

1,515

Resolved to a listing

1,463

Exchange listings

1,987

Articles stored

4,882

Companies (1,515)

Search name or tickr

Exchange

Search

Shenzhen Airport Co Ltd

000089 · CN · hint SZSE

C2:000089

0

no news

Midea Group Co Ltd

000333 · CN · hint SZSE

C2:000333 US:MQCLY ADR

0

no news

Luzhou Lao Jiao Co Ltd

000568 · CN · hint SZSE

C2:000568

0

no news

Spectra AB

0A0L · GB · hint LSE

LN:0A0L US:STKTR US:SCTAY ADR

0

no news

Fox Corporation

0A0X · GB · hint LSE

LN:0A0X US:FOXA US:FOX

2

3d ago

Uber Technologies Inc

0A1U · GB · hint LSE

LN:0A1U US:UBER

34

2d ago

Bloom Energy Corporation

0A4L · GB · hint LSE

LN:0A4L US:BE

24

1d ago

DexCom, Inc.

0A4M · GB · hint LSE

LN:0A4M US:DXCM

3

2d ago

ACADIA Pharmaceuticals Inc.

0A4W · GB · hint LSE

LN:0A4W US:ACAD

4

3d ago

Jack Henry & Associates, Inc.

0A6D · GB · hint LSE

LN:0A6D US:JKHY

15

3d ago

The Cigna Group

0A77 · GB · hint LSE

LN:0A77 US:CIG

6

2d ago

Carvana Co.

0A79 · GB · hint LSE

LN:0A79 US:CVNA

4

3d ago

Palantir Technologies Inc.

GB · hint LSE

LN:0A7R US:PLTR

44

2d ago

Market & macro

Stories that move prices without naming a company — conflict, rates, disasters, policy. From primary wires.

CP Capital Acquires Ascent Athens, a 200-Unit Class A Multifamily Community in Athens, Georgia

BusinessWire · 24h ago

aggregator

Family offices are making a bullish bet on the stock market, according to CNBC Family Office Portfolio Tracker

CNBC · 24h ago

primary wire

This is 'one of the worst' credit card habits, experts say—41% of cardholders do it

CNBC · 1d ago

primary wire

First Financial Corporation and First Illinois Corporation Sign Merger Agreement

GlobalNewsWire · 1d ago

aggregator

Kansas City Fed's Schmid says inflation 'stubborn' and 'sticky,' policy rate not restrictive

CNBC · 1d ago

primary wire

Form 8.3 Picton Property Income Limited

GlobalNewsWire · 1d ago

aggregator

Company news

Nvidia's Unusual Move Answers AI Trade's Key Question; \$40 Trillion US Debt Bomb Ticks; Key Fed Speech Awaited

AAPL · Benzinga · 1d ago

aggregator

ETJ: Very High Tech Exposure Could Be A Risk

AAPL · SeekingAlpha · 1d ago

aggregator

Dollar Tree: The Margin Of Safety Is Finally Gone (Downgrade)

OIC8 · SeekingAlpha · 1d ago

aggregator

VICI Properties: When It Rains Gold, Put Out The Bucket

VICI · SeekingAlpha · 1d ago

aggregator

Realty Income: Business Expansion Has Not Accelerated Growth

O · SeekingAlpha · 1d ago

aggregator

KKR & Co. Inc. Q2: 25% Capital Appreciation Potential From Its Baby Bonds

KKR · SeekingAlpha · 1d ago

aggregator

Research

Hi Insar, ask any financial question

Give me an overview of the market

🔍

Why did the company stock move after its latest earnings...

🔍

What was the company's revenue and EPS for the most...

🔍

Explore what's possible

+ Summarise my portfolio

🔍 Check my strategy fit

🔍 Deep search the news

🔍 Analyse my watchlist

Ask anything

+

🔊

⬆

Every company in the catalogue, with the venues it resolved to. LN:0A1U is the London line, US:UBER the US symbol news is actually fetched with. A 0 with "no news" is a clean zero, not an error — the distinction the run reporting is built around.

Search for stocks, ETFs, and more

My StocksNewsCatalogueGlobal NewsAlerts27My StrategySettings

Global News

The full production catalogue, served by the standalone news component over its HTTP API.

Companies tracked1,515

Resolved to a listing1,463

Exchange listings1,987

Articles stored4,882

Companies (1)

MicrosoftExchangeSearchClear

Microsoft Corp

MSFT · US · no exchange hint

US:MSFT

1172d ago

Market & macro

Stories that move prices without naming a company — conflict, rates, disasters, policy. From primary wires.

CP Capital Acquires Ascent Athens, a 200-Unit Class A Multifamily Community in Athens, Georgia

BusinessWire · 24h ago

aggregator

Family offices are making a bullish bet on the stock market, according to CNBC Family Office Portfolio Tracker

CNBC · 24h ago

primary wire

This is 'one of the worst' credit card habits, experts say—41% of cardholders do it

CNBC · 1d ago

primary wire

First Financial Corporation and First Illinois Corporation Sign Merger Agreement

GlobalNewsWire · 1d ago

aggregator

Kansas City Fed's Schmid says inflation 'stubborn' and 'sticky,' policy rate not restrictive

CNBC · 1d ago

primary wire

Form 8.3 Picton Property Income Limited

GlobalNewsWire · 1d ago

aggregator

Research

Hi Insar, ask any financial question

Give me an overview of the market

Why did the company stock move after its latest earnings...

What was the company's revenue and EPS for the most...

Explore what's possible

+ Summarise my portfolio

Check my strategy fit

Deep search the news

Analyse my watchlist

Ask anything

+🔊⬆

Search by name or ticker. The result shows the resolved listing (US:MSFT) and the stored article count.

Search for stocks, ETFs, and more

My StocksNewsCatalogueGlobal NewsAlerts27My StrategySettings

← Back to Global News

Microsoft Corp

MSFT · US · resolved

US:MSFT

News (117)

NVIDIA CFO Says We Have Invested \$50B In Frontier Labs; OpenAI's Existing, Planned Commitments Represent Around 12GW Of NVIDIA Compute

-Conference Call

Benzinga · 2d ago · also mentions 1 other

IREN Q4 Preview: Stock Down 48% From Highs, Will Data Center Monetization Shine Through in Earnings?

Iren stock is down in 2026, but more data center deals and monetization could help boost earnings and the share price.

Benzinga · 2d ago

Bringing ChatGPT For Teachers To More U.S. School Districts; Data Privacy Agreement Spanning 16 States; Data Privacy Agreement Provides Common Framework For Districts To Evaluate ChatGPT For Teachers Against Student Data Privacy Requirements

https://openai.com/index/bringing-chatgpt-for-teachers-to-more-us-school-districts/

Benzinga · 2d ago

10 Information Technology Stocks With Whale Alerts In Today's Session

This whale alert can help traders discover the next big trading opportunities. Whales are entities with large sums of money and we track their transactions here at Benzinga on our options activity scanner. Traders often

Benzinga · 2d ago · also mentions 7 other

The Fed Faces A Dangerous Choice For Investors; Nvidia Earnings Ahead

Dangerous Choice For Investors Please click here for an enlarged chart of SPDR S&P 500 ETF Trust (NYSE:SPY) which represents the benchmark stock market index S&P 500 (SPX). Note the following: The

Benzinga · 2d ago · also mentions 6 other

Microsoft, Amazon, Google Court Kimi K3 Maker Despite Bessent Blacklist Threat

Microsoft, Amazon and Google are in talks with Kimi K3 maker Moonshot AI over revenue sharing despite Scott Bessent's blacklist threat.

Benzinga · 2d ago · also mentions 3 other

The Government Took Equity In Quantum Stocks, So Should We

Explore quantum computing investing: why profits are years away, why broad/AI funds miss it, and how QTUM vs WQTM ETFs offer diversified exposure– learn...

SeekingAlpha · 2d ago · also mentions 6 other

WQTM: Quantum Computing Offers Promise, But This Product May Not Represent The Most Prudent Proxy

WQTM vs. QTUM: Explore quantum computing ETF outlook, AUM surge, risks of pure-play firms, and why diversified QTUM may fit better. Read the full analysis here.

Research

Hi Insar, ask any financial question

Give me an overview of the market

Why did the company stock move after its latest earnings...

What was the company's revenue and EPS for the most...

Explore what's possible

+ Summarise my portfolio

Check my strategy fit

Deep search the news

Analyse my watchlist

Ask anything

+🔊⬆

One company's feed. "also mentions 1 other" is the shared-article model: a story naming several companies is stored once and linked to each, never duplicated per company.

Read this screenshot critically — it shows a real limitation. Several of those headlines are macro or sector round-ups rather than Microsoft stories. They are there because Benzinga tags a round-up against every ticker it mentions, and **ticker-native attribution is trusted unconditionally**: the provider stated the ticker, so `relevance.ts` short-circuits and the deterministic name check never runs. That is deliberate — the check exists to catch *wrong* attribution, and a provider naming the ticker is not wrong. But **"certain" is not the same as "central"**, and the 99.8% figure in the Results table measures the former.

Ranking a company's feed by how *central* it is to the story is a headline- salience problem, and it is one of the things deliberately not attempted here. The `relevance` score is stored on every link so a consumer can filter on it, and `roundupPenaltySql()` already down-weights stories that name many companies — what is missing is applying that at read time in this view.

The two timing numbers

They measure different things.

A 26-minute run — fetching news for every company, day to day. Finnhub's free plan allows 60 requests per minute, and we ask about one company per request. $1,434 \text{ companies} \div 55 \text{ per minute} \approx 26 \text{ minutes}$. **The code is not slow — it is deliberately waiting** to stay inside the free limit. A paid tier would cut this to a few minutes and change nothing else.

A ~9-minute first start — this happens **once**. On first boot we look up all 1,515 tickers to find their exchanges, via OpenFIGI. Without a key that allows 25 lookups a minute; with a **free** key, 250 a minute and larger batches — hence 9 minutes rather than 38. After the first time it never happens again.

Design decisions worth knowing

- **Provider-agnostic.** Sources sit behind one interface, ordered by config. Adding one is a config change, not a rewrite.
- **Idempotent.** Articles are keyed on a cleaned-up URL, so a daily re-run stores nothing new. Verified: zero duplicate URLs.
- **"No news" is not "provider refused."** They look identical if you only count articles. The first is normal; the second means coverage is dying. They are recorded separately, along with every provider tried per company — because a working fallback can hide a broken primary.
- **No language model.** Article text is written by strangers; letting a model read it and decide what gets stored would hand outsiders influence over the pipeline. The rules used instead are unit-tested against real failures.

Do we need to pay? — and what to buy

Not for coverage. 94.6% is already reachable free and ticker-native.

Yes, for the reading experience. Three problems are not solvable with any free tier, and they are the ones a user actually feels:

1. **Every headline leaves our platform.** The reader lands on Yahoo, Benzinga, Seeking Alpha or Fool — a different site, a different layout, each time.
2. **Some of those sites demand a sign-up or a subscription** before showing the article. We cannot tell in advance which, so the reader plays paywall roulette on links we presented as ours.
3. **We cannot legally fix that ourselves.** Free and low-cost APIs — Finnhub, Marketaux, Alpha Vantage — license a **headline, a short teaser and a link**. Their terms require the link out. Displaying the full text we did not license, or scraping the publisher's page for it, is republishing someone else's copyrighted work. Financial publishers enforce this actively, and it is not a risk worth taking to save a click.

So the requirement "show the article in our own UI" is a licensing decision, not an engineering one. The fix is to buy content we are permitted to display.

Recommendation: Benzinga's licensed newsfeed as primary

Benzinga owns its newsroom, so it can license what an aggregator cannot. Its paid tiers are **explicitly embeddable** — the full body and the image may be published inside our own interface, with no redirect and no third-party sign-in. Its free tier is headline + teaser + link, which is exactly the constraint we are trying to escape.

Requirement	How Benzinga meets it
Show the article without redirecting	Paid tiers license the full body and image for display on our platform
Always fresh, automatically	Pull REST with <code>updatedSince</code> , plus WebSocket, TCP stream and webhooks — push, not polling
No per-company rate ceiling	A firehose of the newsroom, not one request per company — the cost stops scaling with catalogue size
One consistent source	One publisher, one voice, one layout — no paywall roulette

That last row also solves the scale problem in the section above: polling costs one request per company per day, so 50,000 companies is 50,000 requests. A stream costs the same regardless of how many companies we track.

The honest limitation. Benzinga publishes roughly 130–160 full articles a day from a US-focused newsroom. That is deep, not wide — it will not cover Chinese A-shares or thin European lines. It replaces the *quality* of our feed, not all of its *breadth*.

The two-provider architecture

Exactly two, which is the practical minimum that still covers the catalogue:

Tier	Provider	Licence	Reader experience
Primary	Benzinga (paid)	Full body embeddable	Reads in our UI. No redirect, no sign-in
Fallback	Marketaux	Headline + link	Opens the publisher — only for companies Benzinga does not cover

Everything needed to make this switch already exists: providers sit behind one adapter interface and are ordered by `NEWS_PROVIDER_ORDER` , so adding Benzinga is a new adapter plus a config change, and Finnhub can be dropped the day the licensed feed proves out. The article schema already carries `image_url` and a summary; a licensed body needs one nullable `body` column and a UI that renders it instead of linking out.

Pricing is not published — content licensing is quoted per customer, via `licensing@benzinga.com` . The brief offers to provision a paid key, so the ask is a quote for the embeddable-body tier at ~1,500 companies, with room to grow.

Known limits

- **52 companies (3.4%) unresolved** — see the table above; most are foreign secondary listings our identifier sources do not index, not delistings.
- **Name matching is the weakest component.** Some pairs are undecidable from names alone: "Adobe Systems" vs "Adobe" is indistinguishable from "Prudential" vs "Prudential Financial" — the first is one company, the second is two. Those route to a flagged, reduced-confidence path rather than a guess.
- **Some publishers block us.** Bloomberg returns an error to automated requests, so those stories carry no image. We do not disguise the service as a browser to get around it.
- **One ingest worker**; parallel workers need shared rate limiting.

Full detail: [README.md](#) · [COMPARISON.md](#) · [API.md](#) · [OPERATIONS.md](#) · [data/listings-mapping.csv](#)